

Graph Databases

Gábor Szárnyas
CWI, Amsterdam

5 November 2021

Main topics

What are graph databases?


How are they different from relational databases? (e.g. DuckDB, Postgres)

Why are they popular?

How could they work better?

What is a graph database (GDBMS)?

A graph database is a **specialized DBMS** that provides a **graph-oriented data model** and **an API optimized for graph operations**.



graph data model



graph API



graph visualization



Cypher



TigerGraph

GSQL

Oracle Labs
PGX

PGQL



JanusGraph


Gremlin



Amazon Neptune

Gremlin,
SPARQL

Key GDBMS building blocks



property graph
data model



graph query
language



graph
visualization



subgraph
matching

relational
queries

path
queries

stored
procedures

Data model: Property graph



The schema can be optional or mandatory. Lots of decisions and nuances -> [LDBC PGSWG](#)

Query language: Cypher

MATCH

```
(p1:Person)-[:KNOWS]-(p2:Person),  
(p1)<-[:MEMBER]-(f:Forum)-[:MEMBER]->(p2),
```

subgraph
matching

```
(p1)-[:KNOWS*]-(p3:Person)
```

WHERE NOT (f)-[:MEMBER]->(p3)

RETURN p1, f, count(p2), count(p3)




relational
operators

path query
(Kleene-star)

Person

name: Alice

speaks: [en, fr]





KNOWS

MEMBER

since: 2017-05-03

Forum

title: Drums



"p1"	"f"	"count (p2) "	"count (p3) "
{ B }	{ E }	1	1

Query language: Cypher

MATCH

```
(p1:Person)-[:KNOWS]-(p2:Person),  
(p1)<-[:MEMBER]-(f:Forum)-[:MEMBER]->(p2),
```

subgraph
matching

```
(p1)-[:KNOWS*]-(p3:Person)
```

WHERE NOT (f)-[:MEMBER]->(p3)

RETURN p1, f, count(p2), count(p3)

relational
operators


path query
(Kleene-star)

Person

name: Alice

speaks: [en, fr]

KNOWS

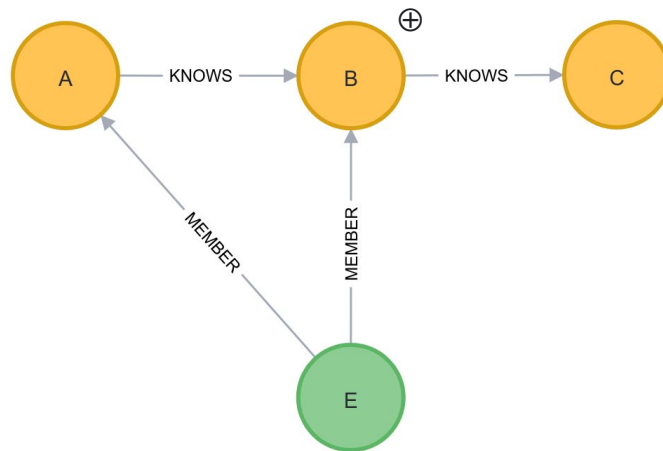


MEMBER

since: 2017-05-03

Forum

title: Drums




Graph Query Language (GQL)

New ISO standard with Cypher-like syntax:

```
USE my_social_graph
MATCH (p:Person)-[:FRIEND*{1,2}]->(friend_or_foaf)
WHERE friend_or_foaf.age > $age AND p.country = $country
RETURN count(*)
```

<https://opencypher.org/articles/2019/09/12/SQL-and-now-GQL/>

<https://ldbouncil.org/event/fourteenth-tuc-meeting/attachments/stefan-plantikow-gql.pdf>



Pattern matching

- **basic graph pattern**
- complex graph pattern
- recommendation query

Subgraph matching (Cypher)



Category: **Basic graph pattern**

MATCH

```
(p1:Person)<-[:MEMBER]-(f:Forum)-[:MEMBER]->(p2:Person),  
(p1)-[:KNOWS]-(p2)
```

WHERE p1.id < p2.id

RETURN p1, p2, f



Results:

(A, B, D)

(A, B, E)

(B, C, E)

Subgraph matching (Cypher)



Category: **Basic graph pattern**

MATCH

```
(p1:Person)<-[:MEMBER]-(f:Forum)-[:MEMBER]->(p2:Person),  
(p1)-[:KNOWS]-(p2)
```

WHERE p1.id < p2.id

RETURN p1, p2, f



Results:

(A, B, D)

(A, B, E)

(B, C, E)

Subgraph matching (Cypher)



Category: **Basic graph pattern**

MATCH

```
(p1:Person)<-[:MEMBER]-(f:Forum)-[:MEMBER]->(p2:Person),  
(p1)-[:KNOWS]-(p2)
```

WHERE p1.id < p2.id

RETURN p1, p2, f



Results:

(A, B, D)

(A, B, E)

(B, C, E)

Subgraph matching (Cypher)



Category: **Basic graph pattern**

MATCH

```
(p1:Person)<-[:MEMBER]-(f:Forum)-[:MEMBER]->(p2:Person),  
(p1)-[:KNOWS]-(p2)
```

WHERE p1.id < p2.id

RETURN p1, p2, f



Results:

(A, B, D)


(A, B, E)

(B, C, E)

Subgraph matching (SQL)

Edge tables: knows(person1id, person2id); member(forumid, personid)

Basic graph pattern: equijoins (SPJ)



Q: $m1 \bowtie m2 \bowtie \text{knows}$

SELECT

m1.personid, m2.personid, m1.forumid

FROM member m1


JOIN member m2

ON m1.forumid = m2.forumid

JOIN knows

ON knows.person1id = m1.personid

AND knows.person2id = m2.personid



Pattern matching

- basic graph pattern
- **complex graph pattern**
- recommendation query

Subgraph matching (Cypher)



Category: **Complex graph pattern**

```
MATCH (f:Forum)-[:MEMBER]->(p1:Person)
```

```
OPTIONAL MATCH (f)-[:MEMBER]->(p2:Person)
```

```
WHERE p1.id < p2.id AND NOT (p1)-[:KNOWS]-(p2)
```

```
RETURN f, count(p2)
```



Results:

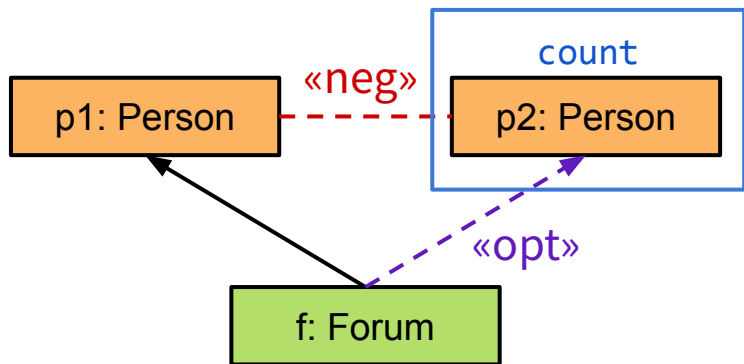
(D, 0)

(E, 1)

Subgraph matching (SQL)

Edge tables: knows(person1id, person2id); member(forumid, personid)

Complex graph pattern: equijoins, outer joins, aggregation (SPOJG)



```
SELECT m1.forumid, count(m2.personid)
```

```
FROM member m1
```

```
LEFT OUTER JOIN member m2
```

```
ON m1.forumid = m2.forumid
```

```
AND m1.personid < m2.personid
```


```
WHERE NOT EXISTS (SELECT true FROM knows
```

```
WHERE person1id = m1.personid
```

```
AND person2id = m2.personid)
```

```
GROUP BY m1.forumid
```

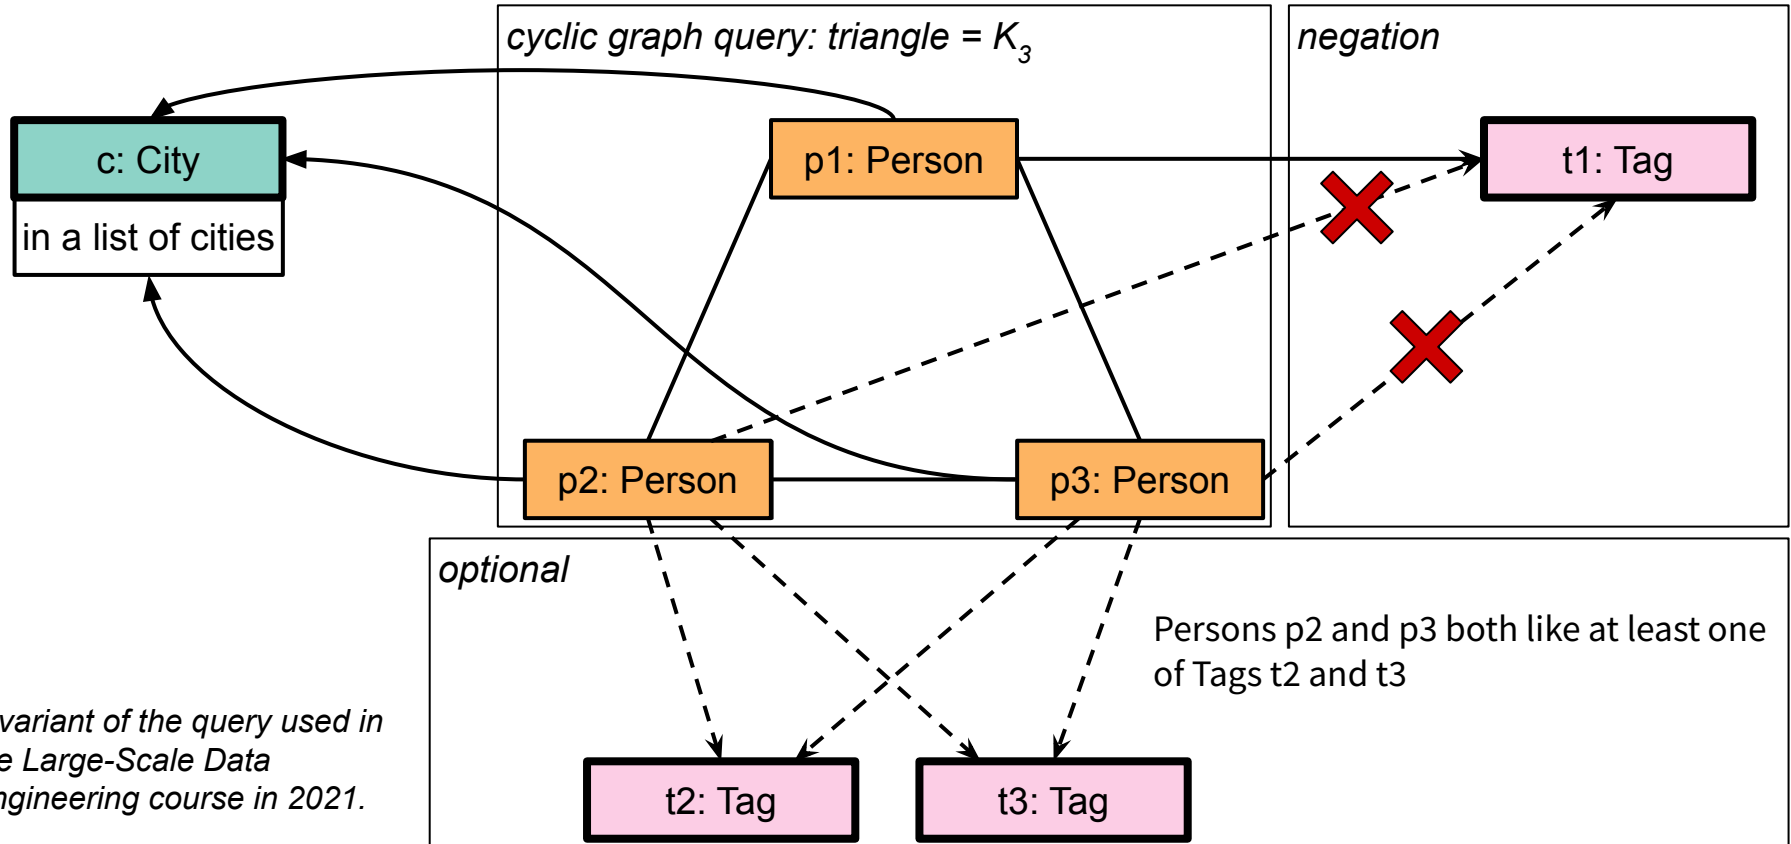
Y (member \bowtie member \triangleright knows)



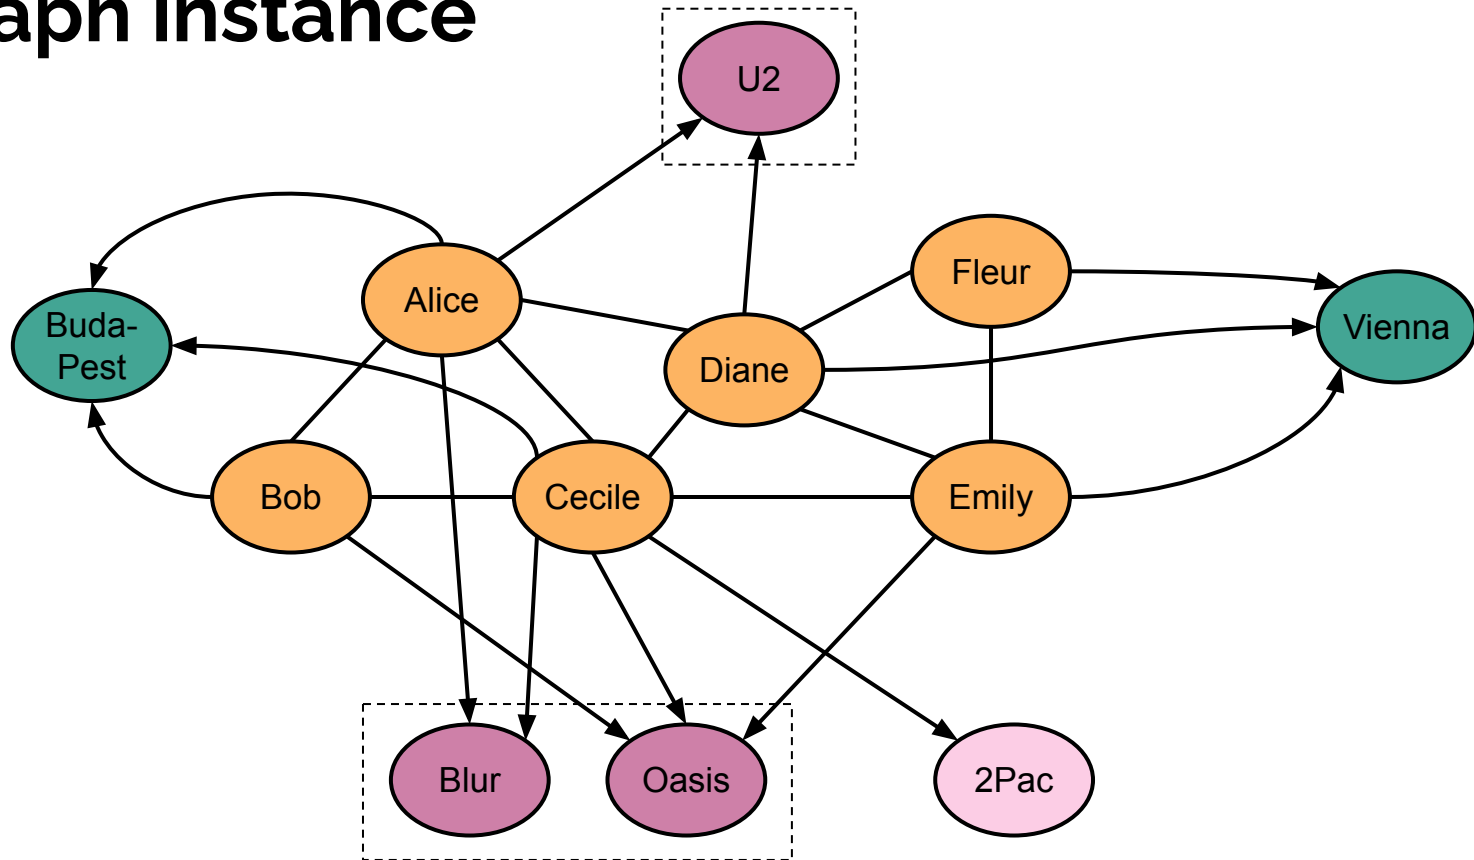
Pattern matching

- basic graph pattern
- complex graph pattern
- **recommendation query**

Recommend concert tickets at a discount



Graph instance



Recommendation query in Cypher

```
MATCH (c:City),
      (p1:Person)-[:LIVES_IN]->(c),
      (p2:Person)-[:LIVES_IN]->(c),
      (p3:Person)-[:LIVES_IN]->(c),
      (p1)-[:KNOWS]-(p2)-[:KNOWS]-(p3)-[:KNOWS]-(p1),
      (p1)-[:LIKES]->(t1:Tag {name: 'U2'}),
      (t2:Tag {name: 'Blur'}),
      (t3:Tag {name: 'Oasis'})
WHERE c.name IN ['Budapest', 'Vienna']
      AND NOT (p2)-[:LIKES]->(t1)
      AND NOT (p3)-[:LIKES]->(t1)
      AND ((p2)-[:LIKES]->(t2) OR (p2)-[:LIKES]->(t3))
      AND ((p3)-[:LIKES]->(t2) OR (p3)-[:LIKES]->(t3))
      AND p1.id < p2.id
      AND p2.id < p3.id
RETURN p1.name, p2.name, p3.name
```

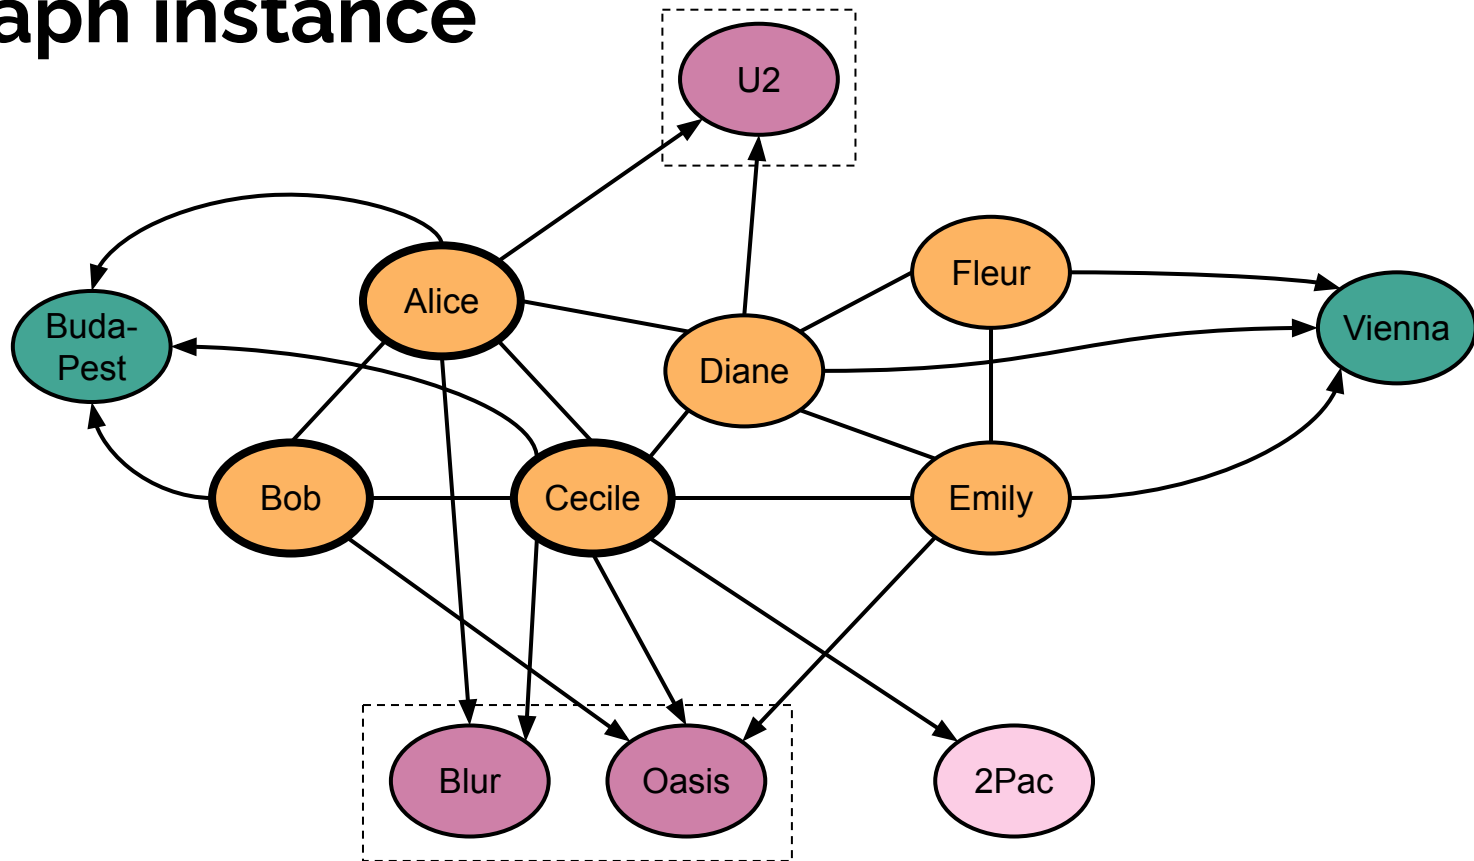
cyclic graph query: triangle = K_3

negation

optional

"p1.name"	"p2.name"	"p3.name"
"Alice"	"Bob"	"Cecile"

Graph instance



Recommendation query in SQL

```
SELECT DISTINCT p1.name, p2.name, p3.name
```

```
FROM City
```

```
JOIN Person p1 ON p1.livesin = City.id
```

```
JOIN Person p2 ON p2.livesin = City.id
```

```
JOIN Person p3 ON p3.livesin = City.id
```

```
JOIN knows k1 ON k1.person1id = p1.id
```

```
JOIN knows k2 ON k2.person1id = p2.id AND k2.person1id = k1.person2id
```

```
JOIN knows k3 ON k3.person1id = p3.id AND k3.person1id = k2.person2id AND k3.person2id = k1.person1id
```

cyclic graph query: triangle = K_3

```
JOIN Tag t1 ON t1.name = 'U2'
```

```
JOIN likes l1 ON l1.personid = p1.id AND l1.tagid = t1.id
```

```
JOIN Tag t2 ON t2.name = 'Blur'
```

```
JOIN Tag t3 ON t3.name = 'Oasis'
```

```
JOIN likes l2 ON l2.personid = p2.id AND (l2.tagid = t2.id OR l2.tagid = t3.id)
```

```
JOIN likes l3 ON l3.personid = p3.id AND (l2.tagid = t2.id OR l2.tagid = t3.id)
```

optional

```
WHERE City.name IN ('Budapest', 'Vienna')
```

```
AND NOT EXISTS (SELECT true FROM likes WHERE likes.personid = p2.id AND likes.tagid = t1.id)
```

```
AND NOT EXISTS (SELECT true FROM likes WHERE likes.personid = p3.id AND likes.tagid = t1.id)
```

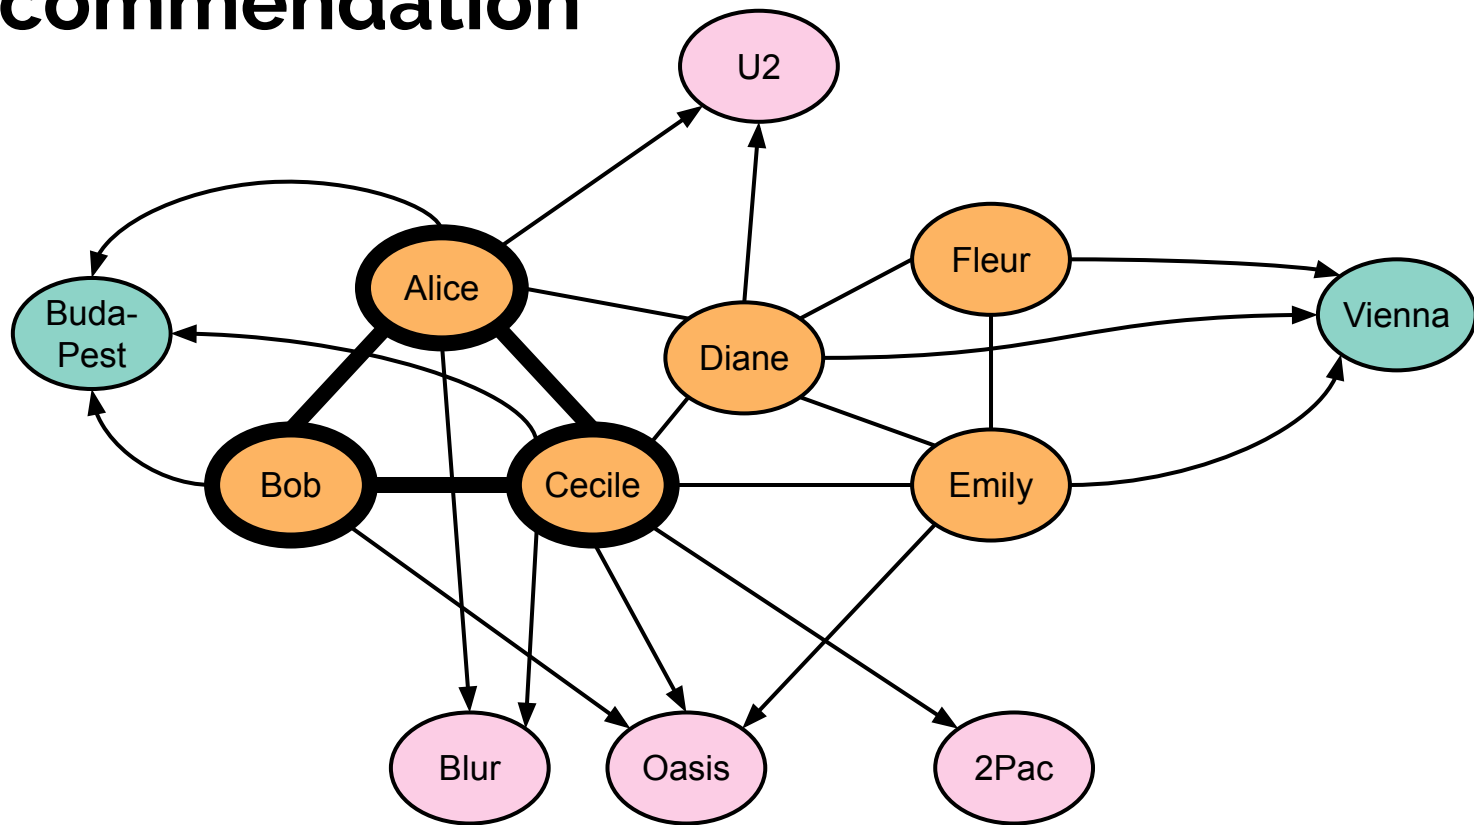
negation

```
AND p1.id < p2.id
```

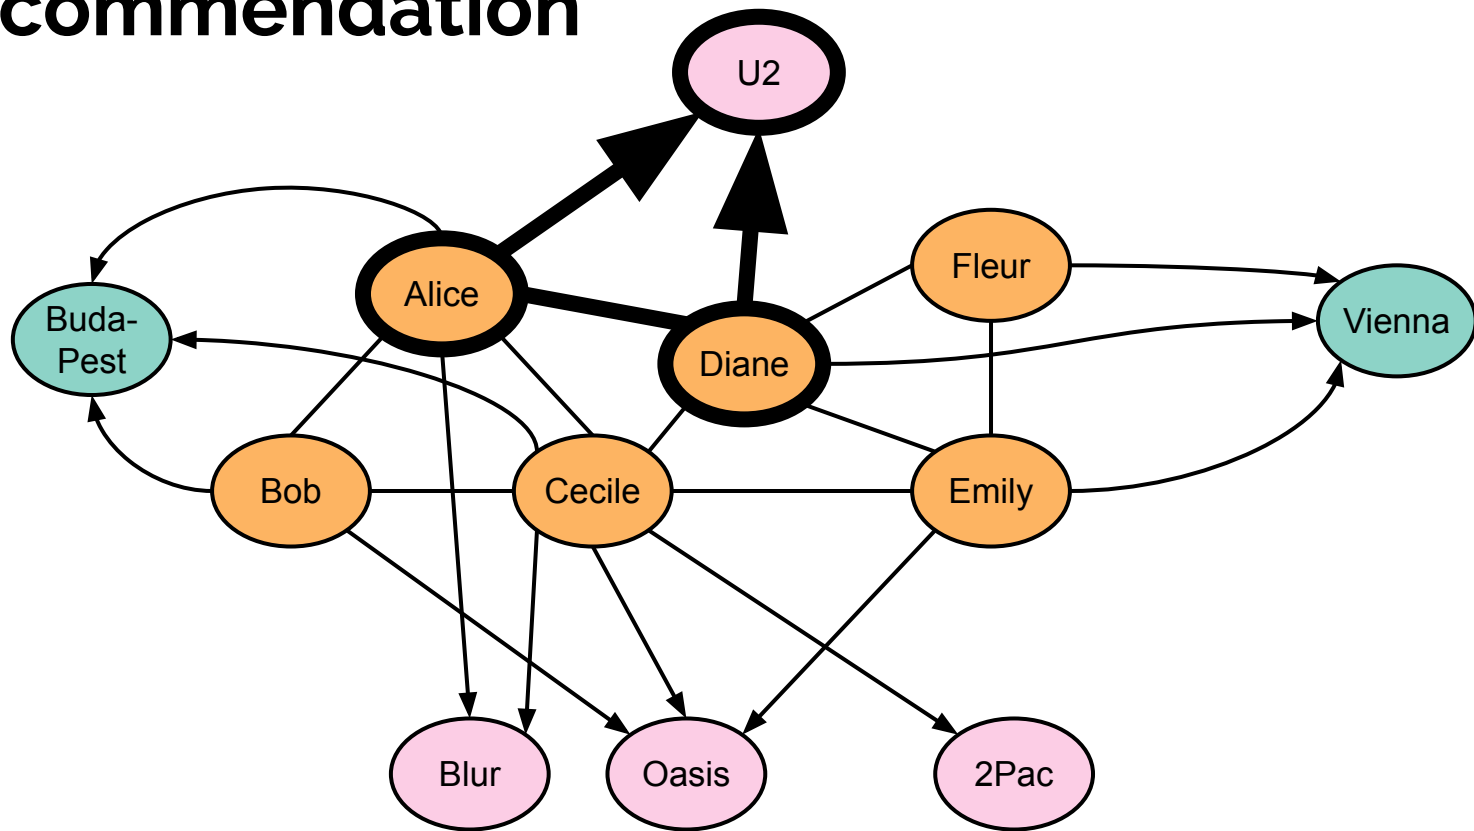
```
AND p2.id < p3.id
```

Cypher query is approx. $\frac{1}{2}$ of the SQL query

Recommendation






Recommendation




Complexity of subgraph matching

Subgraph isomorphism is NP-hard but *matching a given subgraph* is **polynomial**. Still, the complexity of evaluating a **triangle query with binary joins** is **provably suboptimal**, $O(|E|^2)$



Triggered by many-to-many edges and skewed distributions.

Worst-case optimal **multi-way join algorithms** are needed, which have a complexity of just $O(|E|^{1.5})$ for this query.



Research on multiway joins

Subject to research in the last ~15 years:

- **FOCS'08** bounds on complexity
- **PODS'12** Generic-Join algorithm
- **SIGMOD'16** GraphflowDB demo
- **VLDB'19** query optimization in GraphflowDB
- **VLDB'20** integration into Umbra


Industry implementations: RelationalAI, XTDB

Recommendation is a main use case for GDBMSs

[UWaterloo Survey](#) (2017, 2020):

The second most popular application was the use of a graph-based application data to personalize user interactions and provide better recommendations for the customers of a business.

Reason for popularity: ? (maybe visualization?)



Path queries

- **tree query**
- unweighted path query
- weighted shortest path query
- connected components

Tree query: Cypher

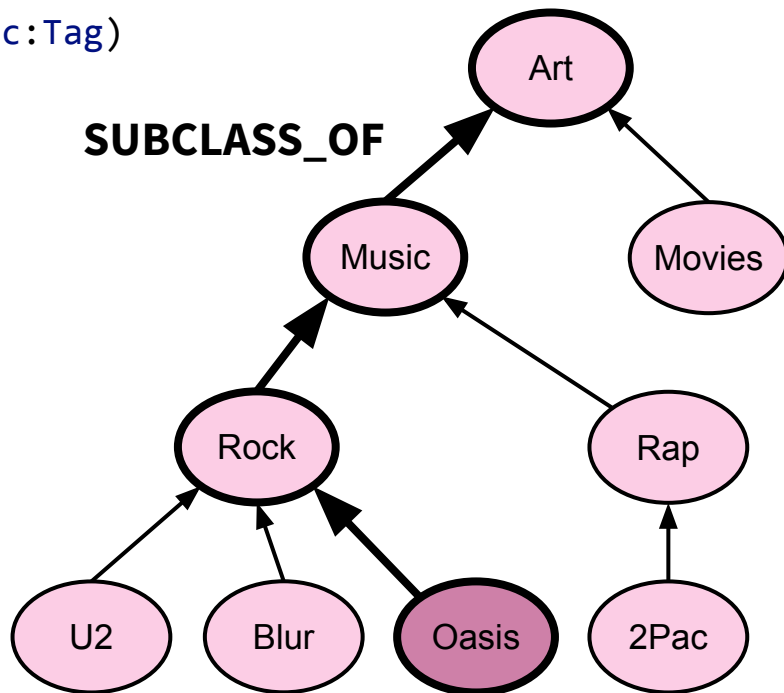
```
MATCH (:Tag {name: 'Oasis'})-[:SUBCLASS_OF*]->(sc:Tag)  
RETURN sc.name AS superclass
```

superclass

Rock

Music

Art



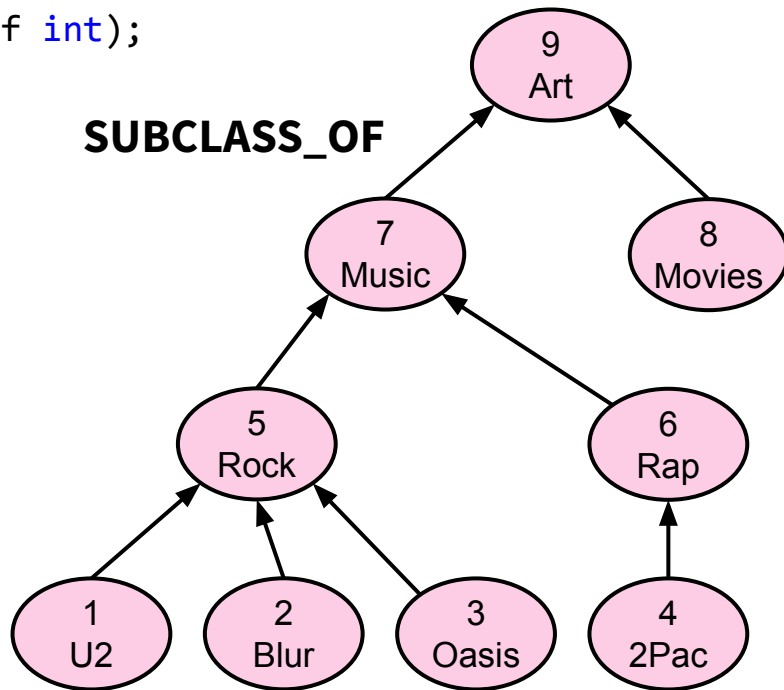
Tree query: SQL data

```
CREATE TABLE tag(id int, name varchar, subclassof int);
```

```
INSERT INTO tag VALUES
```

```
(1, 'U2', 5),  
(2, 'Blur', 5),  
(3, 'Oasis', 5),  
(4, '2Pac', 6),  
(5, 'Rock', 7),  
(6, 'Rap', 7),  
(7, 'Music', 9),  
(8, 'Movies', 9),  
(9, 'Art', NULL)
```

```
;
```



Tree query with string concat

```
WITH RECURSIVE tag_hierarchy(id, source, path) AS (
```

```
SELECT id, name, name AS path
```

```
FROM tag
```

```
WHERE subclassof IS NULL
```

```
UNION ALL
```

```
SELECT tag.id, tag.name, (tag.name || ' -> ' || tag_hierarchy.path)
```

```
FROM tag, tag_hierarchy
```

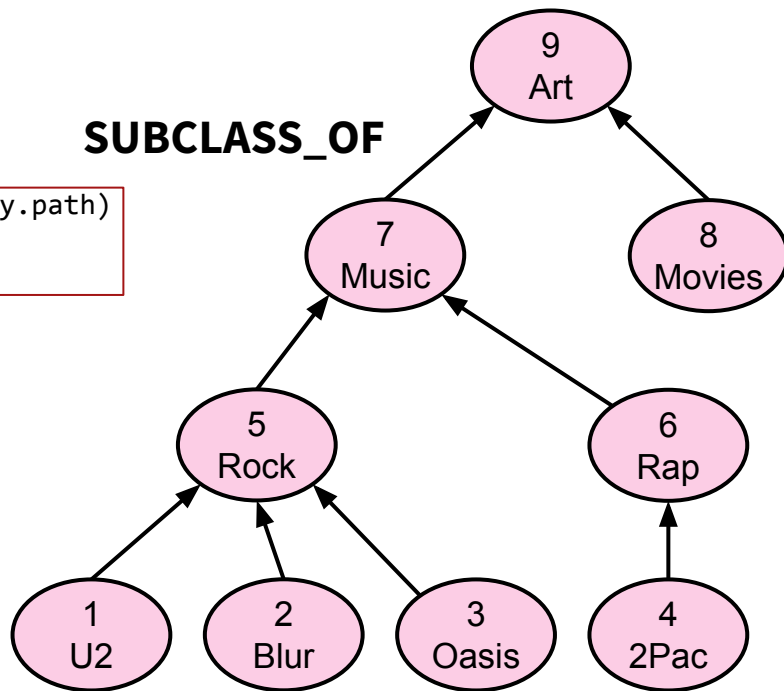
```
WHERE tag.subclassof = tag_hierarchy.id
```

```
)
```

```
SELECT source, path
```

```
FROM tag_hierarchy;
```

source	path
Art	Art
Music	Music -> Art
Movies	Movies -> Art
Rock	Rock -> Music -> Art
Rap	Rap -> Music -> Art
U2	U2 -> Rock -> Music -> Art
Blur	Blur -> Rock -> Music -> Art
Oasis	Oasis -> Rock -> Music -> Art
2Pac	2Pac -> Rap -> Music -> Art

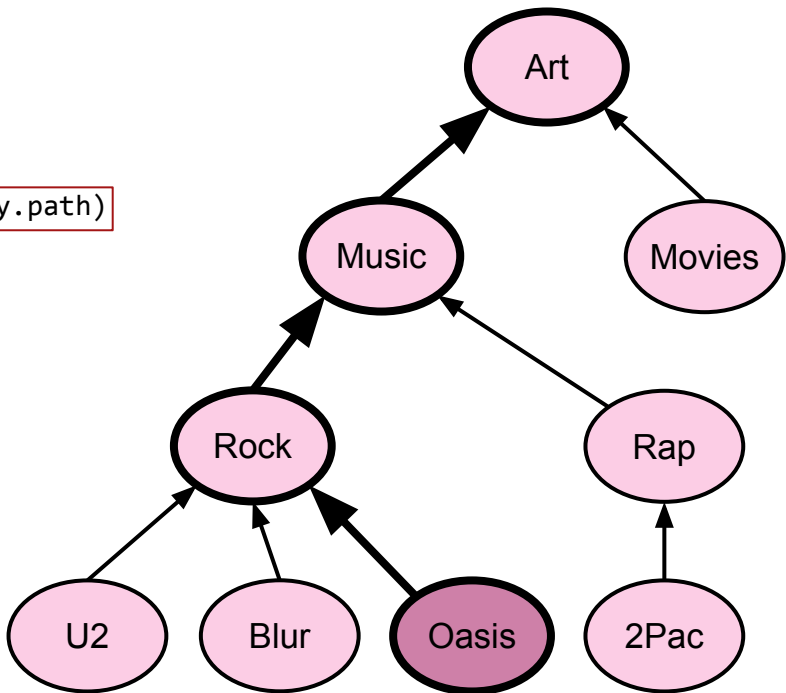



Tree query with list concat

DuckDB 0.3.1-dev
🏆 to Laurens

```
WITH RECURSIVE tag_hierarchy(id, source, path) AS (  
  SELECT id, name, [name] AS path  
  FROM tag  
  WHERE subclassof IS NULL  
  UNION ALL  
  SELECT tag.id, tag.name, list_prepend(tag.name, tag_hierarchy.path)  
  FROM tag, tag_hierarchy  
  WHERE tag.subclassof = tag_hierarchy.id  
)  
SELECT source, unnest(path) AS superclass  
FROM tag_hierarchy  
WHERE source = 'Oasis';
```

source	superclass
Oasis	Oasis
Oasis	Rock
Oasis	Music
Oasis	Art






Path queries

- tree query
- **unweighted path query**
- weighted shortest path query
- connected components

Unweighted shortest path in Cypher

```
MATCH p=shortestPath(  
  (start:Person {name: 'Bob'})-[:KNOWS*]-(end:Person {name: 'Fleur'}))  
RETURN length(p) AS length
```




Result:

length

3

Unw. SP query: Data in SQL

Graphs can be represented in the relational model with PKs and FKs
(primary keys and foreign keys)



Person


id [PK]	name
1	Alice
2	Bob
3	Cecile
4	Diane
5	Emily
6	Fleur

knows

person1id [FK]	person2id [FK]
1	2
1	3
1	4
2	3
3	4
3	5
4	5
4	6
5	6
all edges backwards (optional)	

Unw. SP query: Data in SQL

Graphs can be represented in the relational model with PKs and FKs
(primary keys and foreign keys)



Person

start →

id [PK]	name
1	Alice
2	Bob
3	Cecile
4	Diane
5	Emily
6	Fleur

end →

knows

person1id [FK]	person2id [FK]
1	2
1	3
1	4
2	3
3	4
3	5
4	5
4	6
5	6
all edges backwards (optional)	

Unweighted shortest path query in SQL

```
WITH RECURSIVE paths(startPerson, endPerson, path, level, endPersonReached) AS (  
    SELECT person1id AS startPerson, person2id AS endPerson,  
        [person1id, person2id]::bigint[] AS path, 1 AS level,  
        max(CASE WHEN p2.name = 'Fleur'  
            THEN true ELSE false END) OVER () AS endPersonReached  
    FROM knows  
    JOIN Person p1 ON p1.id = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p1.name = 'Bob'  
UNION ALL  
    SELECT paths.startPerson AS startPerson, person2id AS endPerson,  
        array_append(path, person2id) AS path, level + 1 AS level,  
        max(CASE WHEN p2.name = 'Fleur'  
            THEN true ELSE false END) OVER () AS endPersonReached  
    FROM paths  
    JOIN knows ON paths.endPerson = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p2.id != ALL(paths.path)  
        AND NOT paths.endPersonReached  
SELECT path, level, endPersonReached AS epr  
FROM paths;
```

Unweighted shortest path query in SQL

```
WITH RECURSIVE paths(startPerson, endPerson, path, level, endPersonReached) AS (  
  SELECT person1id AS startPerson, person2id AS endPerson,  
         [person1id, person2id]::bigint[] AS path, 1 AS level,  
         max(CASE WHEN p2.name = 'Fleur'  
                THEN true ELSE false END) OVER () AS endPersonReached  
  FROM knows  
  JOIN Person p1 ON p1.id = knows.person1id  
  JOIN Person p2 ON p2.id = knows.person2id  
  WHERE p1.name = 'Bob'  
  UNION ALL  
  SELECT paths.startPerson AS startPerson, person2id AS endPerson,  
         array_append(path, person2id) AS path, level + 1 AS level,  
         max(CASE WHEN p2.name = 'Fleur'  
                THEN true ELSE false END) OVER () AS endPersonReached  
  FROM paths  
  JOIN knows ON paths.endPerson = knows.person1id  
  JOIN Person p2 ON p2.id = knows.person2id  
  WHERE p2.id != ALL(paths.path)  
         AND NOT paths.endPersonReached  
  SELECT path, level, endPersonReached AS epr  
  FROM paths;
```

initial edge

reached end node?
w/ window function

adding an edge to the path

reached end node?
w/ window function

cycle detection

check reached end node

Unweighted shortest path query in SQL

```
WITH RECURSIVE paths(startPerson, endPerson, path, level, endPersonReached) AS (
```

```
    SELECT person1id AS startPerson, person2id AS endPerson,  
           [person1id, person2id]::bigint[] AS path, 1 AS level,  
           max(CASE WHEN p2.name = 'Fleur'  
                      THEN true ELSE false END) OVER () AS endPersonReached  
    FROM knows  
    JOIN Person p1 ON p1.id = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p1.name = 'Bob'
```

```
UNION ALL
```

```
    SELECT paths.startPerson AS startPerson, person2id AS endPerson,  
           array_append(path, person2id) AS path, level + 1 AS level,  
           max(CASE WHEN p2.name = 'Fleur'  
                      THEN true ELSE false END) OVER () AS endPersonReached  
    FROM paths  
    JOIN knows ON paths.endPerson = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p2.id != ALL(paths.path)  
           AND NOT paths.endPersonReached)
```

```
SELECT path, level, endPersonReached AS epr
```

```
FROM paths;
```

path	level	epr
[2, 3]	1	false
[2, 1]	1	false
[2, 1, 4]	2	false
[2, 3, 1]	2	false
[2, 1, 3]	2	false
[2, 3, 5]	2	false
[2, 3, 4]	2	false
[2, 1, 4, 3]	3	true
[2, 3, 1, 4]	3	true
[2, 3, 5, 4]	3	true
[2, 3, 4, 1]	3	true
[2, 1, 4, 6]	3	true
[2, 1, 3, 5]	3	true
[2, 3, 5, 6]	3	true
[2, 3, 4, 6]	3	true
[2, 1, 4, 5]	3	true
[2, 1, 3, 4]	3	true
[2, 3, 4, 5]	3	true

Unweighted shortest path query in SQL

```
WITH RECURSIVE paths(startPerson, endPerson, path, level, endPersonReached) AS (  
  SELECT person1id AS startPerson, person2id AS endPerson,  
    [person1id, person2id]::bigint[] AS path, 1 AS level,  
    max(CASE WHEN p2.name = 'Fleur'  
      THEN true ELSE false END) OVER () AS endPersonReached  
  FROM knows  
  JOIN Person p1 ON p1.id = knows.person1id  
  JOIN Person p2 ON p2.id = knows.person2id  
  WHERE p1.name = 'Bob'  
UNION ALL  
  SELECT paths.startPerson AS startPerson, person2id AS endPerson,  
    array_append(path, person2id) AS path, level + 1 AS level,  
    max(CASE WHEN p2.name = 'Fleur'  
      THEN true ELSE false END) OVER () AS endPersonReached  
  FROM paths  
  JOIN knows ON paths.endPerson = knows.person1id  
  JOIN Person p2 ON p2.id = knows.person2id  
  WHERE p2.id != ALL(paths.path)  
    AND NOT paths.endPersonReached  
SELECT path, level  
FROM paths  
JOIN Person ON Person.id = paths.endPerson  
WHERE Person.name = 'Fleur';
```

+ unnest + join to get the names

path	level
[2, 1, 4, 6]	3
[2, 3, 5, 6]	3
[2, 3, 4, 6]	3

Unweighted shortest path query in SQL

```
WITH RECURSIVE paths(startPerson, endPerson, path, level, endPersonReached) AS (  
    SELECT person1id AS startPerson, person2id AS endPerson,  
           [person1id, person2id]::bigint[] AS path, 1 AS level,  
           max(CASE WHEN p2.name = 'Fleur'  
                    THEN true ELSE false END) OVER () AS endPersonReached  
    FROM knows  
    JOIN Person p1 ON p1.id = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p1.name = 'Bob'  
UNION ALL  
    SELECT paths.startPerson AS startPerson, person2id AS endPerson,  
           array_append(path, person2id) AS path, level + 1 AS level,  
           max(CASE WHEN p2.name = 'Fleur'  
                    THEN true ELSE false END) OVER () AS endPersonReached  
    FROM paths  
    JOIN knows ON paths.endPerson = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p2.id != ALL(paths.path)  
           AND NOT paths.endPersonReached  
    SELECT path, level  
FROM paths  
JOIN Person ON Person.id = paths.endPerson  
WHERE Person.name = 'Fleur';
```

cycle detection

Unweighted shortest path query in Postgres 14


```
WITH RECURSIVE paths(startPerson, endPerson, path, level, endPersonReached) AS (  
    SELECT person1id AS startPerson, person2id AS endPerson,  
           ARRAY[person1id, person2id]::bigint[] AS path, 1 AS level,  
           max(CASE WHEN p2.name = 'Fleur' THEN 1 ELSE 0 END) OVER () AS endPersonReached  
    FROM knows  
    JOIN Person p1 ON p1.id = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE p1.name = 'Bob'  
UNION ALL  
    SELECT paths.startPerson AS startPerson, person2id AS endPerson,  
           array_append(path, person2id) AS path, level + 1 AS level,  
           max(CASE WHEN p2.name = 'Fleur' THEN 1 ELSE 0 END) OVER () AS endPersonReached  
    FROM paths  
    JOIN knows ON paths.endPerson = knows.person1id  
    JOIN Person p2 ON p2.id = knows.person2id  
    WHERE paths.endPersonReached = 0  
) CYCLE endPerson SET is_cycle TO true DEFAULT false USING cycle_path  
SELECT path, level, endPersonReached AS epr, is_cycle  
FROM paths  
JOIN Person ON Person.id = paths.endPerson  
WHERE Person.name = 'Fleur'  
AND is_cycle = false;
```

cycle detection

path	level	epr	is_cycle
{2,3,5,6}	3	1	f
{2,3,4,6}	3	1	f
{2,1,4,6}	3	1	f
(3 rows)			


Visualize all unweighted shortest paths

```
MATCH p=allShortestPaths(  
    (p1:Person {name: 'Bob'})-[:KNOWS*]-(p2:Person {name: 'Fleur'})  
RETURN p
```



```
[p IN nodes(p) | p.name] AS persons
```

"persons"
["Bob", "Cecile", "Emily", "Fleur"]
["Bob", "Cecile", "Diane", "Fleur"]
["Bob", "Alice", "Diane", "Fleur"]



Path queries

- tree query
- unweighted path query
- **weighted shortest path query**
- connected components

Weighted shortest paths

Difficult. Alternative: stored procedures, e.g. Postgres has [pgrouting](#) and [MADlib](#)

Oracle example from: <http://aprogrammerwrites.eu/?p=1391>

```
WITH paths (node, path, cost, rnk, lev) AS (  
  SELECT a.dst, a.src || ',' || a.dst, a.distance, 1, 1 FROM arcs a  
  WHERE a.src = :SRC  
  UNION ALL  
  SELECT a.dst, p.path || ',' || a.dst, p.cost + a.distance, Rank () OVER (PARTITION BY a.dst ORDER BY p.cost +  
    a.distance), p.lev + 1  
    FROM paths p  
    JOIN arcs a ON a.src = p.node AND p.rnk = 1  
) SEARCH DEPTH FIRST BY node SET line_no  
  CYCLE node SET lp TO '*' DEFAULT ' '  
, paths_ranked AS (  
  SELECT lev, node, path, cost, Rank () OVER (PARTITION BY node ORDER BY cost) rnk_t, lp, line_no  
    FROM paths WHERE rnk = 1)  
SELECT LPad (node, 1 + 2* (lev - 1), '.') node, lev, path, cost, lp  
  FROM paths_ranked  
  WHERE rnk_t = 1  
  ORDER BY line_no
```



Complex query



A relational simulation of Dijkstra's algorithm

Weighted shortest paths

Cypher: No weighted shortest path construct. In Neo4j there's the Graph Data Science lib

```
MATCH (c1:Customer {id: $c1id}), (c2: Customer {id: $c2id})
```

```
CALL gds.shortestPath.dijkstra.stream({
```

call algorithm

```
  nodeProjection: 'Customer',  
  relationshipProjection: 'TRANSFER',  
  sourceNode: c1,  
  targetNode: c2,  
  relationshipWeightProperty: 'amount'  
})
```

```
YIELD path, totalCost
```

```
RETURN path, totalCost
```


- Unweighted shortest path -> query
- Weighted shortest path -> algorithm

Cheapest path (proposal)

Weighted shortest path -> cheapest path.

```
MATCH p AS CHEAPEST (c1 IS customer)-[t IS TRANSFERS COST t.amount]->*(c2 IS customer)
WHERE c1.cid = :customer1Id
      AND c2.cid = :customer2Id
RETURN p
```


(Proposed to Cypher and GQL)



Path queries

- tree query
- unweighted path query
- weighted shortest path query
- **connected components**

Connected components



Can be formulated as a path query but will not be efficient

```
MATCH p=shortestPath((p1:Person)-[:KNOWS*]-(p2:Person))
WHERE p1 <> p2
RETURN p1.id AS node, count(p2.id)+1 AS componentSize
```

The compiler cannot figure out that this is a CC problem -> stored procedures are better

Connected components algorithm in SQL

In-database connected component analysis (ICDE 2020), implemented in Apache HAWQ

```
r.log_exec("setup", 0, """\n
create table ccgraph as
  select v1, v2 from {0}
  union all
  select v2, v1 from {0}
  distributed by (v1);
""").format(dataset))
```

```
roundno = 0
stackA = []
stackB = []
while True:
    roundno += 1
    ccreps = "ccreps{}".format(roundno)
    r_A = 0
    while r_A == 0:
        r_A = random.randint(-2**63, 2**63-1)
    r_B = random.randint(-2**63, 2**63-1)
    stackA.append(r_A)
    stackB.append(r_B)
```

```
r.log_exec("ccreps", roundno, """\n
create table {ccreps} as
  select v1 v,
    least(axplusb({A}, v1, {B}),
           min(axplusb({A}, v2, {B}))) rep
  from ccgraph
  group by v1
  distributed by (v);
""").format(ccreps=ccreps, A=r_A, B=r_B))
```

```
r.log_exec("ccgraph2", roundno, """\n
create table ccgraph2 as
  select r1.rep as v1, v2
  from ccgraph, {} as r1
  where ccgraph.v1 = r1.v
  distributed by (v2);
""").format(ccreps))
r.log_drop("ccgraph")

graphsize = r.log_exec("ccgraph3", roundno, """\n
create table ccgraph3 as
  select distinct v1, r2.rep as v2
  from ccgraph2, {} as r2
  where ccgraph2.v2 = r2.v
    and v1 != r2.rep
  distributed by (v1);
""").format(ccreps))
r.log_drop("ccgraph2")
r.execute("alter table ccgraph3 rename to ccgraph")

if graphsize == 0:
    break
```

ax+b UDF


```
accA = 1
accB = 0

while True:
    roundno -= 1
    (accA, accB) = (r.axplusb(accA, stackA.pop(), 0),
                   r.axplusb(accA, stackB.pop(), accB))
    if roundno == 0:
        break
    ccrepsr = "ccreps{}".format(roundno)
    ccrepsr1 = "ccreps{}".format(roundno+1)
    r.log_exec("result", roundno, """\n
create table tmp as
  select r1.v as v,
    coalesce(r2.rep, axplusb({A}, r1.rep, {B})) as rep
  from {r1} as r1 left outer join
    {r2} as r2
    on (r1.rep=r2.v)
  distributed by (v);
""").format(A=accA, B=accB, r1=ccrepsr, r2=ccrepsr1)
    r.log_drop(ccrepsr)
    r.log_drop(ccrepsr1)
    r.execute("alter table tmp rename to {}".format(ccrepsr))

r.execute("alter table ccreps1 rename to ccrepsresult")
r.log_drop("ccgraph")
```

Fig. 8. Our implementation of Randomised Contraction in Python/SQL.

Index-free adjacency



“Relational databases can’t join efficiently”

It’s a widespread sentiment in the graph DB community that RDBMSs are inefficient.

A slide from the deck
[“Dgraph Series A 2019”](#)
by Manish R. Jain

Pain

- Joins and traversals in SQL are slow.

A [tweet by Frank McSherry](#)
(influential OS/DB researcher)
about *what he heard* at an
industry graph conference.



Frank McSherry
@frankmcsherry

...

Replying to @frankmcsherry

Have now talked with many people who believe that graphDBs are new and neat because RDBMSs obvs use full table scans.

Defining a Graph Database

(Marko Rodriguez, AT&T's graph systems architect, 2010)

A graph database is any storage system that provides index-free adjacency.²³²⁴

²³There is no "official" definition of what makes a database a graph database. The one provided is my definition (respective of the influence of my collaborators in this area). However, hopefully the following argument will convince you that this is a necessary definition. Given that any database can model a graph, such a definition would not provide strict enough bounds to yield a formal concept (i.e. T).

²⁴There is adjacency between the elements of an index, but if the index is not the primary data structure of concern (to the developer), then there is indirect/implicit adjacency, not direct/explicit adjacency. A graph database exposes the graph as an explicit data structure (not an implicit data structure).

Neighbourhood lookup

Efficient neighbourhood lookups (fetch all neighbours of node x) are important for


- traversals (multi-source BFS, bidirectional BFS)
- pattern matching (many WCOJ algorithms intersect neighbourhoods)

Pointers

Lookup is quick when traversing a “many-to-one” edge.



But the expensive traversals go on “many-to-many” edges:



Tradeoffs of pointers

Pointers lead to:

- memory fragmentation
- contention on allocation (synchronization between threads)
- large memory overhead for every node/edge ($\sim 2x$)
- no compression (another $\sim 4x$)
- “tuple-at-a-time” execution with random access at every step

Better solution:

- buffer manager
- ordered tabular storage w/ compression
- vectorized execution (e.g. vectors of 64 tuples, cache misses happen at once)

Columnar store solution

- Everything from an RDBMS
- A graph data structure that supports **analytical queries** and **updates**

Analytical queries benefit from positional indexing (~pointer-based lookups)

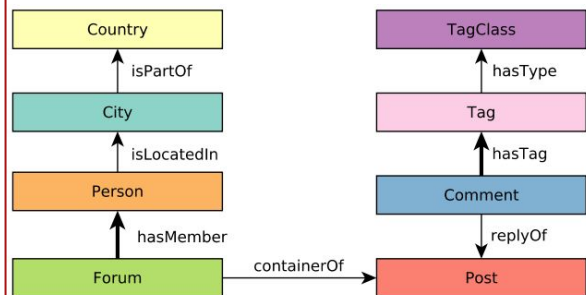
- Invisible join, also supported by DuckDB (Diego Tomé)
- Teseo: packed memory arrays for an updatable CSR-like data structure (Dean de Leo)
- Sortledton: adjacency lists with fast intersection support (Per Fuchs)

Pattern matching performance

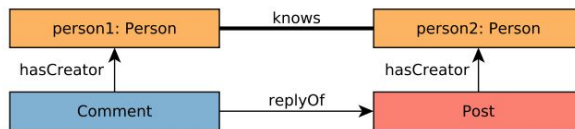
Published at GRADES-NDA 2021, Mhedhbi et al.

Designed to prove that Hannes is wrong

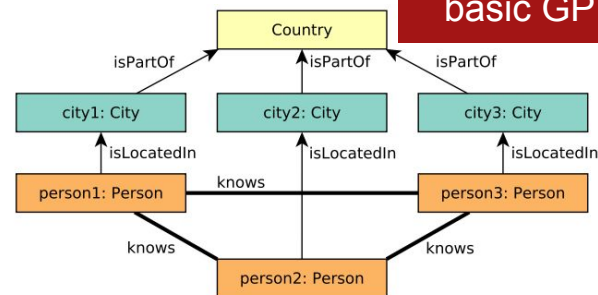
basic GP




(a) Q1.




(b) Q2.




(c) Q3.




(d) Q4.



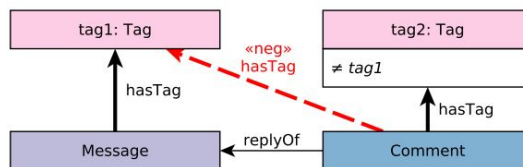
(e) Q5.




(f) Q6.



(g) Q7.



(h) Q8.



(i) Q9.

complex GP

Relational database performance

Best relational systems: columnar stores with compression

- HyPer (industry system, 2010-)
- Umbra: worst-case optimal join support (prototype, 2020-)

Benchmark environment: cloud VM, 370GB RAM, 48 vCPUs




HyPer




UMBRA

GDBMS performance for subgraph queries

- Load the data: 100M vertices, 650M edges
- Run all 9 queries one-by-one
- Environment: cloud VM, 370GB RAM, 48 vCPU cores



Building a GDBMS on an RDBMS



SQL/PGQ: CREATE PROPERTY GRAPH

To be released in 2022


([slides](#))

Property Graph Definition (DDL) - Example

```
CREATE PROPERTY GRAPH aml
  VERTEX TABLES ( account
                  , customer
                  LABEL customer PROPERTIES ( cid, name, city ) )
  EDGE TABLES ( owns SOURCE customers DESTINATION accounts
                 PROPERTIES ( since )
                 , auth_signer SOURCE customer DESTINATION account
                 , transfers
                 SOURCE KEY ( from_id ) REFERENCES accounts ( aid )
                 DESTINATION KEY (to_id) REFERENCES accounts ( aid )
                 LABEL transfers PROPERTIES ( when, amount ) )
```

Explicit label and properties options for customer

Defaults apply for label and all properties.



Columns when and amount are exposed as properties. Columns tid, from_id, and to_id are **not**.



SQL/PGQ: SELECT ... FROM GRAPH_TABLE

To be released in 2022

([slides](#))

Querying PGs – Example 1

New operator* applied to graph (aml), returns table

Retrieve the **info of all customers** who got **more than \$10,000** from customer 100.

```
SELECT gt.cid, gt.name, gt.city, gt.amount
FROM GRAPH_TABLE ( aml,
  MATCH
    ( c1 IS customer ) -[ IS owns ]->
      ( IS account ) -[ t1 IS transfers ]->
        ( IS account ) <-[ IS owns ]- ( c2 IS customer )
  WHERE c1.cid = 100
    AND t1.amount > 10000
  COLUMNS ( c2.cid
            , c2.name
            , c2.city
            , t1.amount )
) gt
```

Edge pattern enclosed in -[]->

Vertex pattern enclosed in ()


COLUMNS defines the shape of the output table. Properties projected out of the MATCH.

* Exact syntax, placement of arguments, etc. currently under discussion/in flux.



GRainDB (UWaterloo)


Neo4j Bloom-style Graph Visualization




G-SQL-style Seamless Table/Graph Querying

```
SELECT DISTINCT Address.zipcode  
FROM (a:vPers)-[:eKnows*1..3]->(b:vPers),  
      Address  
WHERE a.name=Alice AND b.addID=Address.ID
```

Graph Modeling



DuckDB



+ Pre-defined pointer-based joins

+ Factorization

+ Worst-case optimal joins

+ Recursive joins



Summary



Summary

GDBMSs are a new category of DBMSs focused on graph processing

Many languages -> GQL is much needed

Currently, the best approach may be to build on a (columnar) RDBMS

Further readings on graph processing

Talks:

- [GraphflowDB and Modern Query Processing Techniques for GDBMSs](#) (2021) by Semih Salihoglu
- [A Survey of Current Property Graph Query Languages](#) (2020) by Peter Boncz
- [Graph Processing: A Panoramic View and Some Open Problems](#) (2019) by Tamer Ozsu

Survey papers:

- [The future is big graphs: A community view on graph processing systems](#) (2021) by Dagstuhl WG
- [Demystifying Graph Databases: Analysis and Taxonomy of Data Organization, System Designs, and Graph Queries](#) (2020) by Maciej Besta et al.
- [The ubiquity of large graphs and surprising challenges of graph processing: Extended survey](#) (2020) by Siddhartha Sahu et al.
- [Foundations of Modern Graph Query Languages](#) (2017) by Renzo Angles et al.